

Assignment 2

Zahidur Talukder

Student ID- 1001771748

Problem 1

The main requirement of this problem is to implement logistic regression. Logistic regression is a statistical model that in its basic form uses a logistic function to model a binary dependent variable, although many more complex extensions exist. In regression analysis, logistic regression (or logit regression) is estimating the parameters of a logistic model (a form of binary regression).

In this problem, I have used sigmoid function as my activation function and cross entropy for my objective function and finally i performed batch training of my data. Here I will not go for the details theory of the things rather I will discuss about the result of my evaluation.

Here I am attaching the code of my logistic Regression class. Full code is available in my code section.

```
1 ## Code begins here
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5
6 class LogisticRegression:
7
8     def __init__(self):
9         pass
10
11     def sigmoid(self, a):
12         return 1 / (1 + np.exp(-a))
13
14     def train(self, X, y_true, n_iters, learning_rate):
15         n_samples, n_features = X.shape
16         self.weights = np.zeros((n_features, 1))
17         self.bias = 0
18         costs = []
19         for i in range(n_iters):
20             y_predict = self.sigmoid(np.dot(X, self.weights) + self.bias
21 )
22             cost = (- 1 / n_samples) * np.sum(y_true * np.log(y_predict)
23 + (1 - y_true) * (np.log(1 - y_predict)))
24             dw = (1 / n_samples) * np.dot(X.T, (y_predict - y_true))
25             db = (1 / n_samples) * np.sum(y_predict - y_true)
26             self.weights = self.weights - learning_rate * dw
27             self.bias = self.bias - learning_rate * db
28             costs.append(cost)
29             if i % 100 == 0:
30                 print(f"Cost after iteration {i}: {cost}")
31
32         return self.weights, self.bias, costs
33
34     def predict(self, X):
35         y_predict = self.sigmoid(np.dot(X, self.weights) + self.bias)
36         y_predict_labels = [1 if elem > 0.5 else 0 for elem in y_predict
37 ]
38         return np.array(y_predict_labels)[: , np.newaxis]
39
40     def confusion_metrics(self, labels, predictions, threshold):
41         true_positive=0;
```


```

39     false_positive=0;
40     true_negative=0;
41     false_negative=0;
42     for i in range(len(labels)):
43         if labels[i]==1:
44             if predictions[i]>=threshold:
45                 true_positive+=1;
46             else:
47                 false_negative+=1;
48         else:
49             if predictions[i]>=threshold:
50                 false_positive+=1;
51             else:
52                 true_negative+=1;
53     tpf=true_positive/(true_positive + false_negative);
54     fpf=false_positive/(false_positive + true_negative);
55     return tpf,fpf;
56
57 def results(self,labels,predictions):
58     TPF=[];
59     FPF=[];
60     THRESHOLD=[];
61     i=0;
62     #increemental step size for threshold
63     dx_step=0.0002;
64     while(i<=1):
65         threshold=i;
66         tpf,fpf=confusion_metrics(labels,predictions,threshold);
67         TPF.append(tpf);
68         FPF.append(fpf);
69         THRESHOLD.append(threshold);
70         i+=dx_step;
71
72     plt.plot(FPF,TPF);
73     plt.plot(THRESHOLD,THRESHOLD,'--')
74     plt.xlabel("False Positive fraction (FPF)--->")
75     plt.ylabel("True Positive fraction (TPF)--->")
76     plt.title("ROC Curve")
77     plt.show()
78     area = np.trapz(TPF, dx=dx_step)
79     print("AUC:Area under the ROC curve is", area)
80     return;
81

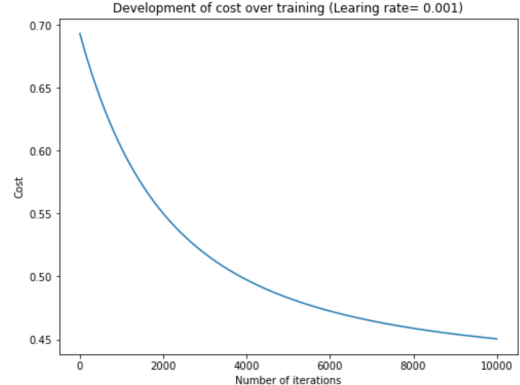
```

Result Evaluation


Here I have a train function for training the model and predict function for predicting the model based on training in the LogisticRegression Class. During the training period I can control the learning rate. For different learning rate the cost function graph is given below:




(a) Cost function for learning rate .0001



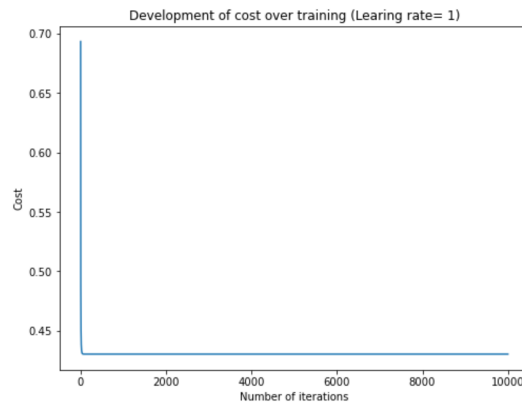
(b) Cost function for learning rate .001



(c) Cost function for learning rate .01



(d) Cost function for learning rate .1



(e) Cost function for learning rate 1


Figure 1: Cost function for different learning rate

Here from the graph we can see that when the learning rate is small, model needs long time to converge and this is obvious. When the learning rate is 0.0001 system doesn't even converge with 10000 iterations, when its 0.001, its about to converge, when it is 0.01, it converges after about 3000 iterations, when 0.1, it converges after 500 iterations and finally when 1, it converges after 100 iterations.

Here blue line indicate the ROC graph and from the figure we can see that the Area Under Curve (AUC) is 0.810019. From the graph we can also see that the accuracy that we get from our training and testing data set is 82.2 and 83.2 percent respectively. This is the most accuracy that we have got from our data.

Since logistic regression is a binary regression it performs well for classification problem. Since our data set is from two multivariate normal distribution which has many common merging area so our model could not perform more than 83 percent accurate. But it performs very good for other types of binary data.

train accuracy: 82.2%
test accuracy: 83.2%



AUC:Area under the ROC curve is 0.810019

Figure 2: ROC Curve with AUC and Accuracy value

Problem 1.2

The main challenge in this problem was to Implement multi-class logistic regression using a soft-max function and cross entropy in MNIST dataset. MNIST-dataset is 60000 training and 10000 test samples that contain images of hand-written numbers in 28x28 pixels. But since we need to take only first 5(0 to 4) classes, the number of training and testing samples was reduced about half. In this experiment I used hot encoding of my target data and train the model. Here I attached the soft-max regression class and complete code is available in my code section.

```
1 ## Code begins here
2
3 from tensorflow.keras.datasets import mnist
4 import pandas as pd
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 class SoftmaxRegression:
9     def __init__(self):
10         pass
11
12     def initialize_param(self,d):
13         np.random.seed(1)
14         params = {}
15         params['w'] = np.random.randn(d,5)/np.sqrt(d)
16         params['b'] = np.zeros((5,1))
17         return params
18
19     def softmax(self,Z):
20         expZ = np.exp(Z - np.max(Z))
21         return expZ / expZ.sum(axis=0, keepdims=True)
22
23     def forward(self, params, X):
24         w = params['w']
25         b = params['b']
26         Z = np.dot(w.T,X) + b
27         A = self.softmax(Z)
28         return A
29
30     def compute_cost(self,A,Y):
31         m = Y.shape[1]
32         cost = (-1/m)*np.sum(Y * np.log(A + 1e-8))
33         return cost
34
35     def backprop(self,X, Y, A):
36         m = Y.shape[1]
37         dw = (1/m) * np.dot(X, (A - Y).T)
38         db = (1/m) * np.sum(A - Y)
39         return dw, db
40
41     def optimise(self, params, X, Y, num_iterations, l_rate):
42         costs = []
43         for i in range(num_iterations):
44             A = self.forward(params, X)
45             cost = self.compute_cost(A, Y)
46             dw, db = self.backprop(X, Y, A)
47
48             params['w'] = params['w'] - l_rate * dw
49             params['b'] = params['b'] - l_rate * db
50
```

```

51         if i % 100 == 0:
52             print("Cost after iteration %i : %f " %(i, cost))
53
54         costs.append(cost)
55     return params, costs
56
57     def predict(self, params, X, Y):
58         w = params['w']
59         #print(w.shape, X.shape)
60         probs = self.forward(params, X)
61         y_hat = np.argmax(probs, axis=0)
62         Y = np.argmax(Y, axis=0)
63         conf_matrix = self.compute_confusion_matrix(Y, y_hat)
64         print("label precision recall")
65         for label in range(5):
66             print(f"{label:5d} {self.precision(label, conf_matrix):9.3f} {self.recall(label, conf_matrix):6.3f}")
67
68         accuracy = self.accuracy(conf_matrix)
69         print("Test Accuracy: ", accuracy)
70         return conf_matrix
71
72     def model(self, d, X_train, Y_train, num_interation, l_rate):
73         params = self.initialize_param(d)
74         params, costs = self.optimise(params, X_train, Y_train, num_interation, l_rate)
75         return params, costs
76
77     def compute_confusion_matrix(self, true, pred):
78         K = len(np.unique(true)) # Number of classes
79         result = np.zeros((K, K))
80         for i in range(len(true)):
81             result[true[i]][pred[i]] += 1
82         return result
83
84     def precision(self, label, confusion_matrix):
85         col = confusion_matrix[:, label]
86         return confusion_matrix[label, label] / col.sum()
87
88     def recall(self, label, confusion_matrix):
89         row = confusion_matrix[label, :]
90         return confusion_matrix[label, label] / row.sum()
91
92     def accuracy(self, confusion_matrix):
93         diagonal_sum = confusion_matrix.trace()
94         sum_of_all_elements = confusion_matrix.sum()
95         return diagonal_sum / sum_of_all_elements
96


```

Result evaluation

The main challenge here was to get the first 5 class of data from MNIST data set. There are 60000 training data and there are 10000 testing data in the dataset for 10 classes. So when I take only 5 classes the number of dataset was about half as we can see in the figure. Again the MNIST dataset is 28×28 pixels image. I convert them to a linear array of $28 \times 28 = 784$. Again I have done hot encoding to my target data as you can see the shape of the data.

Train Input Shape: (784, 30596)
Train Target Shape: (5, 30596)
Test Input Shape: (784, 5139)
Test Target Shape: (5, 5139)

Figure 3: For 5 class of MNIST dataset



(a) Cost Vs Iteration

label	precision	recall
0	0.988	0.994
1	0.989	0.991
2	0.965	0.945
3	0.963	0.972
4	0.978	0.981

Test Accuracy: 0.9766491535318156

(b) Accuracy, Precision and Recall

Figure 4: Result of the Model for Training and Testing

In this problem I have used the softmax function for activation function as you can see in the code and cross entropy for objective function. The model is trained with train data set and you can see the cost vs iteration of my training data. Although the model converges much earlier I iterated it 10000 times as requirement. Then by checking the test data, The accuracy of the model is 97.66 percent.

For finding the precision and recall of my model I defined the function compute confusion matrix for getting confusion matrix .And from that confusion matrix using the function precision and recall I found out the Precision and Recall for each class level as you can see in the figure, the average precision was about 97 percent and recall was 99 percent.

Problem 3

The main challenge here was to implement an Artificial Neural Network(ANN). I have write the code for making the ANN. Here I have used the same MNIST dataset as required. Here I have used both the Sigmoid and ReLu function as activation function. The class for ANN is given here full code is availbale in the code section.

```
1 ## Code begins here
2
3 from tensorflow.keras.datasets import mnist
4 import pandas as pd
5 import numpy as np
6 import matplotlib.pyplot as plt
7
8 class ANN_perceptron:
9     def __init__(self):
10         pass
11
12
13     def sigmoid(self, Z):
14         A = 1/(1+np.exp(-Z))
15         return A
16
17     def relu(self, Z):
18         A = np.maximum(0,Z)
19         cache = Z
20         return A
21
22     def deep_initialize_parameters(self, layer_dims):
23         np.random.seed(1)
24         parameters = {}
25         L = len(layer_dims)                # number of layers in the network
26
27         for l in range(1, L):
28             parameters['W' + str(l)] = np.random.randn(layer_dims[l],
29 layer_dims[l-1])/ np.sqrt(layer_dims[l-1])# *0.01
30             parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))
31         return parameters
32
33     def linear_activation_forward(self, A_prev, W, b, activation):
34         Z = np.dot(W,A_prev) + b
35         linear_cache = (A_prev, W, b)
36
37         if activation == "sigmoid":
38             A = sigmoid(Z)
39         elif activation == "relu":
40             A = relu(Z)
41
42         cache = (linear_cache, Z)
43         return A, cache
44
45     def deep_model_forward(self, X, parameters):
46         caches = []
47         A = X
48         L = len(parameters) // 2
49         for l in range(1, L):
50             A_prev = A
51             A, cache = self.linear_activation_forward(A_prev, parameters
52 ['W' + str(l)], parameters['b' + str(l)], activation = "relu")
```

```

51         caches.append(cache)
52         AL, cache = self.linear_activation_forward(A, parameters['W' +
53 str(L)], parameters['b' + str(L)], activation = "sigmoid")
54         caches.append(cache)
55         return AL, caches
56
57     def compute_cost(self, AL, Y):
58         m = Y.shape[1]
59         cost = (1./m) * (-np.dot(Y,np.log(AL).T) - np.dot(1-Y, np.log(1-
60 AL).T))
61         cost = np.squeeze(cost)
62         assert(cost.shape == ())
63         return cost
64
65     def linear_backward(self, dZ, cache):
66         A_prev, W, b = cache
67         m = A_prev.shape[1]
68
69         dW = 1./m * np.dot(dZ,A_prev.T)
70         db = 1./m * np.sum(dZ, axis = 1, keepdims = True)
71         dA_prev = np.dot(W.T,dZ)
72         return dA_prev, dW, db
73
74     def relu_backward(self, dA, cache):
75         Z = cache
76         dZ = np.array(dA, copy=True)
77         dZ[Z <= 0] = 0
78         assert (dZ.shape == Z.shape)
79         return dZ
80
81     def sigmoid_backward(self, dA, cache):
82         Z = cache
83         s = 1/(1+np.exp(-Z))
84         dZ = dA * s * (1-s)
85         assert (dZ.shape == Z.shape)
86         return dZ
87
88     def linear_activation_backward(self, dA, cache, activation):
89         linear_cache, activation_cache = cache
90         if activation == "relu":
91             dZ = self.relu_backward(dA, activation_cache)
92             dA_prev, dW, db = linear_backward(dZ, linear_cache)
93
94         elif activation == "sigmoid":
95             dZ = self.sigmoid_backward(dA, activation_cache)
96             dA_prev, dW, db = linear_backward(dZ, linear_cache)
97
98         return dA_prev, dW, db
99
100     def deep_model_backward(self,AL, Y, caches):
101         grads = {}
102         L = len(caches) # the number of layers
103         m = AL.shape[1]
104         Y = Y.reshape(AL.shape)
105
106         dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))
107
108         current_cache = caches[L-1]
109         grads["dA" + str(L)], grads["dW" + str(L)], grads["db" + str(L)]
110 = self.linear_activation_backward(dAL, current_cache, activation = "

```


```

sigmoid")
108     for l in reversed(range(L-1)):
109         # lth layer: (RELU -> LINEAR) gradients.
110         current_cache = caches[l]
111         dA_prev_temp, dW_temp, db_temp = self.
linear_activation_backward(grads["dA" + str(l + 2)], current_cache,
activation = "relu")
112         grads["dA" + str(l + 1)] = dA_prev_temp
113         grads["dW" + str(l + 1)] = dW_temp
114         grads["db" + str(l + 1)] = db_temp
115
116     return grads
117
118     def update_parameters(self, parameters, grads, learning_rate):
119         L = len(parameters) // 2
120         for l in range(L):
121             parameters["W" + str(l+1)] = parameters["W" + str(l+1)] -
learning_rate * grads["dW" + str(l+1)]
122             parameters["b" + str(l+1)] = parameters["b" + str(l+1)] -
learning_rate * grads["db" + str(l+1)]
123
124         return parameters
125
126     def deep_layer_model(self, X, Y, layers_dims, learning_rate=0.0075,
num_iterations=3000): #lr was 0.009
127         costs = []
128         parameters = self.deep_initialize_parameters(layers_dims)
129         for i in range(0, num_iterations):
130             AL, caches = self.deep_model_forward(X, parameters)
131             cost = self.compute_cost(AL, Y)
132             grads = self.deep_model_backward(AL, Y, caches)
133             parameters = self.update_parameters(parameters, grads,
learning_rate)
134             costs.append(cost)
135             if i % 100 == 0:
136                 print ("Cost after iteration %i: %f" % (i, cost))
137         return parameters, costs
138
139
140     def predict(self, params, X, Y):
141         probs, caches = self.deep_model_forward(X, parameters)
142         y_hat = np.argmax(probs, axis=0)
143         Y = np.argmax(Y, axis=0)
144         conf_matrix=self.compute_confusion_matrix(Y, y_hat)
145         print("label precision recall")
146         for label in range(5):
147             print(f"{label:5d} {self.precision(label, conf_matrix):9.3f}
{self.recall(label, conf_matrix):6.3f}")
148
149         accuracy = self.accuracy(conf_matrix)
150         print("Test Accuracy: ", accuracy)
151         return conf_matrix
152
153     def compute_confusion_matrix(self, true, pred):
154         K = len(np.unique(true)) # Number of classes
155         result = np.zeros((K, K))
156         for i in range(len(true)):
157             result[true[i]][pred[i]] += 1
158         return result
159

```

```
160     def precision(self, label, confusion_matrix):
161         col = confusion_matrix[:, label]
162         return confusion_matrix[label, label] / col.sum()
163
164     def recall(self, label, confusion_matrix):
165         row = confusion_matrix[label, :]
166         return confusion_matrix[label, label] / row.sum()
167
168     def accuracy(self, confusion_matrix):
169         diagonal_sum = confusion_matrix.trace()
170         sum_of_all_elements = confusion_matrix.sum()
171         return diagonal_sum / sum_of_all_elements
172
```

Result evaluation



(a) Cost Vs Iteration


label	precision	recall
0	0.992	0.996
1	0.996	0.994
2	0.979	0.972
3	0.988	0.989
4	0.987	0.991

Test Accuracy: 0.9883245767659078

(b) Accuracy, Precision and Recall

Figure 5: Result of the Model for Training and Testing for ReLU activation function

Here I have used 3 hidden layer with nodes 256, 64,32. These are dense layer and for the first case I have used ReLU function as my activation function. For this case as we can see from the figure I have done 10000 iteration and the precision and recall for each level is given above. Again the model has accuracy of 98.83 percent for test data set. So in this case the system performs better than Multi-class logistic regression.



(a) Cost Vs Iteration

label	precision	recall
0	0.997	0.995
1	0.996	0.996
2	0.984	0.986
3	0.995	0.993
4	0.994	0.994

Test Accuracy: 0.9929947460595446

(b) Accuracy, Precision and Recall

Figure 6: Result of the Model for Training and Testing for Sigmoid activation function

As you can see from the figure, for the Second case I have used Sigmoid function as my activation function. For this case as we can see from the figure I have done 10000 iteration and the precision and recall for each level is given above. Again the model has accuracy of 99.29 percent for test data set. So in this case the system performs better than Multi-class logistic regression as well.

So for my case the Artificial Neural Network(ANN) performs better than Multi-class logistic regression. The main advantage of ANNs over logistic regression models lies in their hidden layers of nodes. In fact, a special ANN with no hidden node has been shown to be identical to a logistic regression model. That was obvious for my case also. And since the data set was large enough, there was no overfitting. Thus in both the cases the model performs really well.